

van Emde Boas Trees and a Faster Shortest Path Algorithm

Tomer Azuz

Vrije Universiteit Amsterdam

25/08/2022

Overview

- 1 Introduction
- 2 Construction
- 3 Optimized Shortest Path Algorithm

Overview

- 1 Introduction
- 2 Construction
- 3 Optimized Shortest Path Algorithm

- A data structure that supports all dynamic-set operations in $O(\log \log u)$ worst-case time
 - *insert, delete, lookup, successor, predecessor, min, max*
- Matches a lower bound for the predecessor problem

Overview

- 1 Introduction
- 2 Construction
- 3 Optimized Shortest Path Algorithm

Bit Vector

- Maintain a vector of n bits
- Elements in the set correspond to the indices in which the bit is set

0	0	1	0	0	1	0	1
0	1	2	3	4	5	6	7

A bit vector representing the set $\{2, 5, 7\}$

Bit Vector - $\text{successor}(x)$


- Starting from index $x + 1$, scan all bits until a 1 is encountered
- $O(n)$ worst-case running time

0	0	0	0	0	0	0	1
0	1	2	3	4	5	6	7

Calling $\text{successor}(0)$ on the singleton set $\{7\}$


Speeding-up *successor*

- Augment a binary tree on top of the bit vector
- Bits are stored in the leaves
- Each non-leaf node stores the logical-or of its two children




A binary tree augmented on top of a bit vector representing the set $\{2, 5, 7\}$

Successor of 2



Start from index 2, scan the tree bottom-up until a node is accessed from the left and its right child contains a 1




Head down towards the leaves, always taking the leftmost node containing a 1

successor: Time-Complexity

- Height of a binary tree is $\log n$
- Two scans through at most $\log n$ levels are required
- *successor* worst-case running time: $O(\log n)$

Augment a tree of degree \sqrt{u} on top of the bit vector


- u : Range of all possible values stored in the set
- Assume u is an even power of two, such that \sqrt{u} is an integer
- Middle level is an array of \sqrt{u} bits, where each entry i indicates whether cluster i is empty or not
- Bottom level is divided into \sqrt{u} clusters of size \sqrt{u}




A tree of degree \sqrt{u} augmented on top of a bit vector representing the set $\{1, 8, 11, 13, 14\}$

- $high(x)$: Determines the cluster number of x by computing x/\sqrt{u}
- $low(x)$: Determines the index of x within its cluster by computing $x \bmod \sqrt{u}$
- $index(x, y)$: Constructs an element from its cluster number x and its position y within its cluster by computing $x \cdot \sqrt{u} + y$

Successor of 1




Successor of 1



Find the next non-empty cluster by scanning the *summary* array

Successor of 1



Scan the next non-empty cluster.
The first position containing a 1 is the successor of 1

successor(x): Time Complexity

In the worst case:


- All \sqrt{u} bits in x 's cluster are scanned
- All \sqrt{u} bits in the *summary* array are scanned
- All \sqrt{u} bits in the next non-empty cluster are scanned

Worst-case running time of *successor*: $O(\sqrt{u})$

van Emde Boas Tree Structure

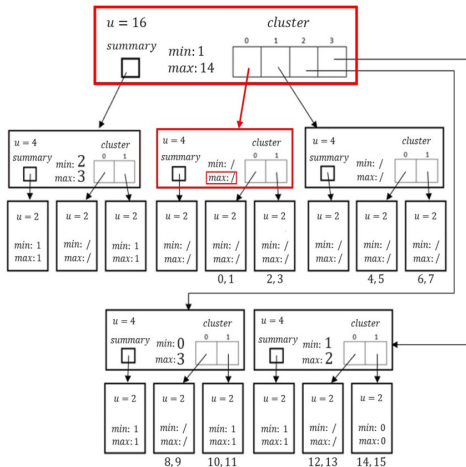
- Each non-leaf vEB node maintains the following attributes:
 - u : Range of all possible values in the set
 - min and max : Store the minimum and maximum elements
 - $summary$: A pointer to a vEB node of size $\lceil \sqrt{u} \rceil$
 - $cluster$: Array of pointers to $\lceil \sqrt{u} \rceil$ vEB nodes of size $\lfloor \sqrt{u} \rfloor$
 - A leaf vEB node only maintains the attributes u , min and max
- Recursive structure
 - Starting from a root node of size u , at each level of the tree u decreases by a factor of square root down to a base size of 2
 - The height of the resulting tree is $\log \log u$

van Emde Boas Tree



A van Emde Boas tree representing the set $\{1, 8, 11, 13, 14\}$

Successor of 1



successor(x)

if V is a leaf:

```
if x == 0 and max == 1:
```

```

return 1

```

```
else
```

```
return null
```

```
if V is nonempty and  $x < \min$ :
```

```
return min
```

```
maxLow = clusters[high(x)].max
```

```
if maxLow != null and low(x) < maxLow:
```

```
offset = cluster[high(x)].successor(low(x))
```

```
return index(high(x), offset)
```

else

```
succCluster = summary.successor(high(x))
```

```
if succCluster == null:
```

```
return null
```

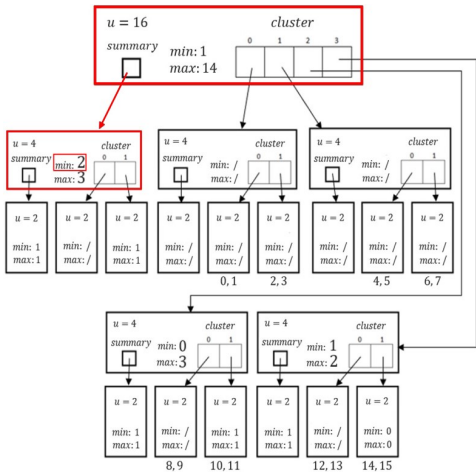
```
else
```

```
offset = cluster[succCluster].min
```

```
return index(succCluster, offset)
```

Search within 1's cluster

Successor of 1

 $\text{successor}(x)$

if V is a leaf:

```
if x == 0 and max == 1:
```

```
return 1
```

else

```
return null
```

```
if V is nonempty and  $x < \min$ :
```

```
return min
```

```
maxLow = clusters[high(x)].max
```

```
if maxLow != null and low(x) < maxLow:
```

```
offset = cluster[high(x)].successor(low(x))
```

```
return index(high(x), offset)
```

else

```
succCluster = summary.successor(high(x))
```

```
if succCluster == null:
```

```
return null
```

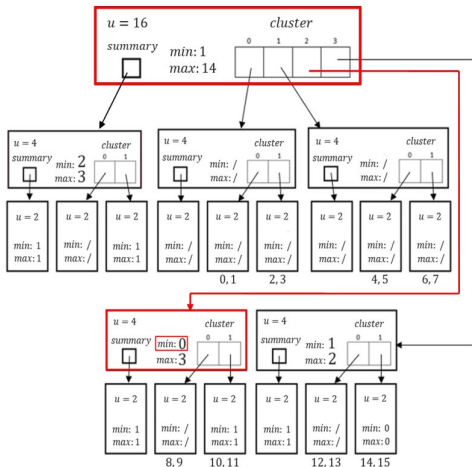
else

```
offset = cluster[succCluster].min
```

```
return index(succCluster, offset)
```

Query the *summary* for the next non-empty cluster

Successor of 1

 $\text{successor}(x)$

if V is a leaf:

```
if x == 0 and max == 1:
```

```
return 1
```

else

```
return null
```

```

if V is nonempty and  $x < \min$ :

```

```

return min

```

```
maxLow = clusters[high(x)].max
```

```
if maxLow != null and low(x) < maxLow:
```

```
offset = cluster[high(x)].successor(low(x))
```

```
return index(high(x), offset)
```

else

```
succCluster = summary.successor(high(x))
```

```
if succCluster == null:
```

```
return null
```

else

```
offset = cluster[succCluster].min
return index(succCluster, offset)
```

The minimum element in the next non-empty cluster is the successor of 1

successor: Time Complexity

- Two recursive calls
 - Only one is performed - $O(\log \log u)$
 - Remaining operations have $O(1)$ cost

```
successor(x)
  if V is a leaf:
    if x == 0 and max == 1:
      return 1
    else
      return null
  if V is nonempty and x < min:
    return min
  maxLow = clusters[high(x)].max
  if maxLow != null and low(x) < maxLow:
    offset = cluster[high(x)].successor(low(x))
    return index(high(x), offset)
  else
    succCluster = summary.successor(high(x))
    if succCluster == null:
      return null
    else
      offset = cluster[succCluster].min
      return index(succCluster, offset)
```

Overview

- 1 Introduction
- 2 Construction
- 3 Optimized Shortest Path Algorithm**

Dijkstra's Algorithm Optimization

- Shortest Path Problem: Find a path with the lowest weight between two vertices in a graph
- Dijkstra's Algorithm Optimization: Use a van Emde Boas tree as a priority queue

	Time Complexity	Space Complexity
Array	$O(V ^2)$	$O(V)$
Min-Heap	$O(V + E \cdot \log V)$	$O(V)$
Fibonacci Heap	$O(V \cdot \log V + E)$	$O(V)$
vEB	$O((V + E) \cdot \log \log c)$	$O(V + c)$

Dijkstra's Algorithm Complexity Analysis

Thank you!